

1. A basic and often useful data structure is a doubly linked list. In internal memory it uses  $O(N)$  space and supports the following operations: Given a pointer to an element  $x$  in the list, a new element can be inserted before or after  $x$  in  $O(1)$  time, and  $x$  can be deleted in  $O(1)$  time. Similarly, given a pointer to  $x$ , the  $K$ -element sublist starting with  $x$  can be traversed in  $O(K)$  time, for any  $K \geq 1$ . Here your goal is to construct an I/O-efficient doubly-linked list. The structure should occupy  $O(N/B)$  blocks, should support insertions and deletions as above using  $O(1)$  I/Os, and should allow the traversal of any  $K$ -element sublist using  $O(1 + K/B)$  I/Os. Argue that the structure you develop achieves these space and I/O bounds.
2. Let  $X[0..n-1]$  and  $Y[0..n-1]$  be two arrays, each storing a set of  $n$  numbers. Let  $Z[0..n-1]$  be another array, in which each entry  $Z[i]$  has three fields:  $Z[i].x$ ,  $Z[i].y$  and  $Z[i].sum$ . The fields  $Z[i].x$  and  $Z[i].y$  contain integers in the range  $\{0, \dots, n-1\}$ ; the fields  $Z[i].sum$  are initially empty. We wish to store in each field  $Z[i].sum$  the value  $X[Z[i].x] + Y[Z[i].y]$ . A simple algorithm for this is as follows.  
 ComputeSums( $X, Y, Z$ )  
 1: for  $i = 0$  to  $n-1$  do  
 2:    $Z[i].sum = X[Z[i].x] + Y[Z[i].y]$   
 3: end for  
 (i) Analyze the number of I/Os performed by ComputeSums.  
 (ii) Give an algorithm to compute all  $Z[i].sum$  that performs only  $O(\text{Sort}(n))$  I/Os.
3. Prove that in Van Emde boas layout of  $n$  numbers, searching a number needs at most  $4 \log_B(N/B)$  I/Os. If we search all numbers existing in the tree in an arbitrary order, prove a search needs  $2 \log_B(N/B)$  I/Os in average.
4. An  $N$ -element B-tree occupies  $O(N/B)$  disk blocks and supports insertions, deletions, and FIND operations using  $O(\log_B N)$  I/Os. Here your goal is to use the buffering idea to obtain a structure with the same space requirements and the same cost for FIND operations, but which uses only  $O((1/\sqrt{B}) \log_B N)$  amortized I/Os per insertion or deletion. Argue that your structure achieves these space and I/O bounds. For simplicity, you can assume there is not any deletion.
5. Batched range reporting is the following problem: Given a set  $R$  of rectangles in the plane and a set  $P$  of points in the plane, report all pairs  $(p, r)$  such that  $p \in P, r \in R$ , and  $p \in r$ . If  $N = |P| + |R|$  and  $K$  is the number of such pairs  $(p, r)$  reported, your task is to design an algorithm that solves this problem using  $O(\text{sort}(N) + K/B)$  I/Os. Provide a description, analysis, and proof of correctness of your algorithm.