


Massive Data Algorithmics

Lecture 6: Interval Trees

Interval Management

- **Problem:**

- Maintain N intervals with unique endpoints dynamically such that stabbing query with point x can be answered efficiently




- As in (one-dimensional) B-tree case we are interested in

- $O(N/B)$ space
- $O(\log_B N)$ update
- $O(\log_B N + T/B)$


Interval Management: Static Solution

- **Sweep** from left to right maintaining persistent B-tree
 - Insert interval when left endpoint is reached
 - Delete interval when right endpoint is reached




- Query x answered by reporting all intervals in B-tree at time x
 - $O(N/B)$ space
 - $O(\log_B N)$ update
 - $O(\log_B N + T/B)$

Internal Interval Trees




- Base tree on endpoints slab X_v associated with each node v
- Interval stored in highest node v where it contains midpoint of X_v
- Intervals I_v associated with v stored in
 - Left slab list sorted by left endpoint (search tree)
 - Right slab list sorted by right endpoint (search tree)
- Linear space and $O(\log n)$ update

Internal Interval Trees




- Query with x on left side of midpoint of X_{root}
 - Search left slab list left-right until finding non-stabbed interval
 - Recurse in left child
- $\Rightarrow O(\log N + T)$ query bound

Externalizing Interval Tree



- **Natural idea:**
 - Block tree
 - Use B-tree for slab lists
- Number of stabbed intervals in large slab list may be small (or zero)
 - We can be forced to do I/O in each of $O(\log N)$ nodes


Externalizing Interval Tree



Idea:

- Decrease fan-out to $\Theta(\sqrt{B}) \Rightarrow$ height remains $O(\log_B N)$
- $\Theta(\sqrt{B})$ slabs define $\Theta(B)$ multislabs
- Interval stored in two slab lists (as before) and one multislab list
- Intervals in small multislab lists collected in underflow structure
- Query answered in v by looking at 2 slab lists and not $O(\log B)$


Externalizing Interval Tree



Idea:

- Decrease fan-out to $\Theta(\sqrt{B}) \Rightarrow$ height remains $O(\log_B N)$
- $\Theta(\sqrt{B})$ slabs define $\Theta(B)$ **multislabs**
- Interval stored in two slab lists (as before) and one **multislab list**
- Intervals in small multislab lists collected in **underflow structure**
- Query answered in v by looking at 2 slab lists and not $O(\log B)$

Externalizing Interval Tree




• Idea:

- Decrease fan-out to $\Theta(\sqrt{B}) \Rightarrow$ height remains $O(\log_B N)$
- $\Theta(\sqrt{B})$ slabs define $\Theta(B)$ **multislabs**
- Interval stored in two slab lists (as before) and one **multislab list**
- Intervals in small multislab lists collected in **underflow structure**
- Query answered in v by looking at 2 slab lists and not $O(\log B)$


Externalizing Interval Tree

- Base tree: Weight-balanced B-tree with branching parameter $1/4\sqrt{B}$ and leaf parameter B on endpoints
 - Interval stored in highest node v where it contains slab boundary
- Each internal node v contains:
 - Left slab list for each of $\Theta(\sqrt{B})$ slabs
 - Right slab list for each of $\Theta(\sqrt{B})$ slabs
 - $\Theta(B)$ multislab lists
- Interval in set I_v of intervals associated with v stored in
 - Left slab list of slab containing left endpoint
 - Right slab list of slab containing right endpoint
 - Widest multislab list it spans
- If $< B$ intervals in multislab list they are instead stored in underflow structure (\Rightarrow contains $= B^2$ intervals)



Externalizing Interval Tree

- Base tree: **Weight-balanced B-tree** with branching parameter $1/4\sqrt{B}$ and leaf parameter B on endpoints
 - Interval stored in highest node v where it contains slab boundary
- Each internal node v contains:
 - Left slab list** for each of $\Theta(\sqrt{B})$ slabs
 - Right slab list** for each of $\Theta(\sqrt{B})$ slabs
 - $\Theta(B)$ **multislab lists**
- Interval in set I_v of intervals associated with v stored in
 - Left slab list** of slab containing **left endpoint**
 - Right slab list** of slab containing **right endpoint**
 - Widest multislab** list it spans
- If $< B$ intervals in multislab list they are instead stored in **underflow structure** (\Rightarrow contains $= B^2$ intervals)




Externalizing Interval Tree

- Each leaf contains $< B/2$ intervals (unique endpoint assumption)
 - Stored in one block
- Slab lists implemented using B-trees
 - $O(1 + T_v/B)$ query
 - Linear space
 - * We may waste a block for each of the $\Theta(\sqrt{B})$ lists in node
 - * But only $\Theta(\frac{N}{B\sqrt{B}})$ internal nodes
- Underflow structure implemented using static structure
 - $O(\log_B B^2 + T_v/B) = O(1 + T_v/B)$ query
 - Linear space
- Linear space

Externalizing Interval Tree

- **Query** with x
 - Search down tree for x while in node v reporting **all intervals in I_v stabbed by x**
- In node v
 - Query two slab lists
 - **Report** all intervals in **relevant multislab** lists
 - Query **underflow structure**
- Analysis:
 - Visit $O(\log_B N)$ nodes
 - Query slab lists $O(1 + T_v/B)$
 - Query multislab lists $O(1 + T_v/B)$
 - Query underflow structure $O(1 + T_v/B)$


$$\Rightarrow O(\log_B N + T_v/B)$$



Externalizing Interval Tree

- Update ignoring base tree update/rebalancing:

- Search for relevant node: $O(\log_B N)$
- Update two slab lists: $O(\log_B N)$
- Update **multislab list** or **underflow structure**



- Update of **underflow structure** in $O(1)$ I/Os amortized:

- Maintain update block with $\leq B$ updates
- Check of update block adds $O(1)$ I/Os to query bound
- **Rebuild structure** when **B updates** have been collected using $O(B^2/B \log_B B^2) = O(B)$ I/Os (Global Rebuilding)

- \Rightarrow Update in $O(\log_B N)$ I/Os amortized

Externalizing Interval Tree

- Note:

- Insert may increase number of intervals in underflow structure for some multislab to B
- Delete may decrease number of intervals in multislab to B

⇒ Need to move B intervals to/from multislab/underflow structure

- We only move

- Intervals from multislab list when decreasing to size $B/2$
- Intervals to multislab list when increasing to size B

⇒ $O(1)$ I/Os amortized used to move intervals


Base Tree Update

- Before inserting new interval we insert new endpoints in base tree using $O(\log_B N/B)$ I/Os

- Leads to **rebalancing** using splits

⇒ Boundary in v becomes boundary in $\text{parent}(v)$



⇒ Intervals need to be moved




- Move intervals (update secondary structures) in $O(w(v))$ I/Os
 - ⇒ $O(1)$ amortized split bound (weight balanced B-tree)
 - ⇒ $O(\log_B N/B)$ amortized insert bound

Splitting Interval Tree Node

- When v splits we may need to move $O(w(v))$ intervals
 - Intervals in v containing boundary
 - Intervals in $\text{parent}(v)$ with endpoints in X_v containing boundary
- Intervals move to two new slab and **multislab** lists in $\text{parent}(v)$




Splitting Interval Tree Node



- Moving intervals in v in $O(w(v))$ I/Os
 - Collected in left order (and remove) by scanning left **slab lists**
 - Collected in right order (and remove) by scanning right **slab lists**
 - Removed **multislabs** containing boundary
 - Remove from **underflow structure** by rebuilding it
 - Construct lists and **underflow** structure for v and v similarly

Splitting Interval Tree Node




- Moving intervals in $\text{parent}(v)$ in $O(w(v))$ I/Os
 - Collect in left order by scanning left **slab list**
 - Collect in right order by scanning right **slab list**
 - Merge with intervals collected in $v \Rightarrow$ two new **slab lists**
 - Construct new **multislabs** by splitting relevant **multislabs**
 - Insert intervals in small **multislabs** in **underflow structure**

Splitting Interval Tree Node

- Split in $O(1)$ I/Os amortized
 - Space: $O(N/B)$
 - Query: $O(\log_B N/B + T/B)$
 - Insert: $O(\log_B N/B)$ I/Os amortized
- Deletes in $O(\log_B N/B)$ I/Os amortized using global rebuilding:
 - Delete interval as previously using $O(\log_B N/B)$ I/Os
 - Mark relevant endpoint as deleted
 - Rebuild structure in $O(\log_B N/B)$ after $N/2$ deletes
- Note: Deletes can also be handled using fuse operations

Summary/Conclusion: Interval Management

- Interval management corresponds to simple form of $2d$ range search
 - Diagonal corner queries
- We obtained the same bounds as for the $1d$ case
 - Space: $O(N/B)$
 - Query: $O(\log_B N/B + T/B)$
 - Update: $O(\log_B N/B)$



Summary/Conclusion: Interval Management

- Main problem in designing structure:
 - Binary \rightarrow large fan-out
- Large fan-out resulted in the need for
 - Multislabs and multislabs lists
 - Underflow structure to avoid $O(B)$ -cost in each node
- General solution techniques:
 - **Filtering**: Charge part of query cost to output
 - **Bootstrapping**:
 - * Use $O(B^2)$ size structure in each internal node
 - * Constructed using persistence
 - * Dynamic using global rebuilding
 - **Weight-balanced B-tree**: Split/fuse in amortized $O(1)$