



Massive Data Algorithmics

Lecture 9: Algorithms for trees

Graphs





- Massive graphs
 - Web modeling: web crawling
 - Geographic information systems: Modeling terrains by graphs
- Representing graphs
 - Adjacency list
 - Unordered collection of edges



- Massive graphs
 - Web modeling: web crawling
 - Geographic information systems: Modeling terrains by graphs
- Representing graphs
 - Adjacency list
 - Unordered collection of edges


- Key difficulties in designing I/O-efficient graph algorithms
 - Nodes visited in **unpredictable order**.
unstructured access to adjacency lists seems to need at least one I/O per node.
 - Remembering **settled nodes** requires extra data structures-algorithmic changes.





- Many results, **many open questions**.
- **Undirected** case often easier than **directed** cases.
- **Dense** graphs often easier than **sparse** graphs
- Special graph classes often easier
- General Methods: Time-forward processing, PRAM simulation, Graph reduction, ...
- Efficient solutions: MST, CC, Listranking, ...
- Still difficult: BFS, DFS, Shortest paths, ...

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$

- Make rooted
- Preorder ranking
- Postorder ranking
- Computing depth

- Can be simply done with $O(|V|)$ I/Os

- Can be done in $O(\text{sort}(|V|))$?

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$

- Make rooted
- Preorder ranking
- Postorder ranking
- Computing depth

- Can be simply done with $O(|V|)$ I/Os

- Can be done in $O(\text{sort}(|V|))$?

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$

- Make rooted
- Preorder ranking
- Postorder ranking
- Computing depth

- Can be simply done with $O(|V|)$ I/Os


- Can be done in $O(\text{sort}(|V|))$?

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$
 - Make rooted
 - Preorder ranking
 - Postorder ranking
 - Computing depth
- Can be simply done with $O(|V|)$ I/Os
- Can be done in $O(\text{sort}(|V|))$?

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$
 - Make rooted
 - Preorder ranking
 - Postorder ranking
 - Computing depth
- Can be simply done with $O(|V|)$ I/Os
- Can be done in $O(\text{sort}(|V|))$?

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$
 - Make rooted
 - Preorder ranking
 - Postorder ranking
 - Computing depth
- Can be simply done with $O(|V|)$ I/Os
- Can be done in $O(\text{sort}(|V|))$?

Fundamental algorithms for trees




- Fundamental algorithms on tree $T = (V, E)$
 - Make rooted
 - Preorder ranking
 - Postorder ranking
 - Computing depth
- Can be simply done with $O(|V|)$ I/Os
- Can be done in $O(\text{sort}(|V|))$?

Listranking



- Given a link list L , compute for every element of L its distance from the head of L .
- More General:** each element v associated with $w(v)$. Compute $\rho(v)$ where $\rho(v) = \rho(\text{pred}(v)) \oplus w(v)$.

Listranking



- Naive algorithms


Procedure NAÏVELISTRANKING

```
1:  $v \leftarrow h$ 
2:  $\rho \leftarrow 0_l$ 
3: while  $v \neq \text{nil}$  do
4:    $\rho \leftarrow \rho \oplus \omega(v)$ 
5:    $\rho(v) \leftarrow \rho$ 
6:    $v \leftarrow \text{succ}(v)$ 
7: end while
```


$\{0_l \text{ is the left-neutral element w.r.t. } \oplus.\}$

- $O(|V|)$ I/Os with LRU paging strategy

Listranking



- Maintained information for each node
 - Node id
 - Successor id
 - $w(v)$ (known) and $\rho(v)$ (to be computed)
 - extra data depending on applications



- Overall strategy
 - If L fits into memory, load L to the memory.
 - Construct L' with size $2/3|L|$ with removing a large independent set I .
 - Updates the weight of elements in L' so that their weight ranks in L and L' are the same.
 - Recurse on L'
 - Compute the weight rank of elements in I by adding their weights to the weight ranks of their predecessors
- $\mathcal{I}(\mathcal{N}) = \mathcal{O}(\text{sort}(\mathcal{N}))$


$$\mathcal{I}(N) = \begin{cases} \mathcal{O}(\text{scan}(N)) & \text{if } N \leq M \\ \mathcal{I}(\frac{2}{3}N) + \mathcal{O}(\text{sort}(N)) & \text{if } N > M \end{cases}$$

Listranking


Procedure FASTLISTRANKING

- 1: **if** $|L| \leq M$ **then**
- 2: Load list L into main memory, and use procedure NAÏVELISTRANKING to compute the ranks of all elements in L .
- 3: **else**
- 4: Find an independent set I of size at least $N/3$ in L .
- 5: **for** all $v \in L \setminus I$ **do**
- 6: $\text{succ}_{L'}(v) \leftarrow \text{succ}_L(v)$
- 7: $\rho_{L'}(v) \leftarrow \rho_L(v)$
- 8: **end for**
- 9: **for** all $v \in I$ **do**
- 10: **if** $\text{succ}_L(v) \neq \text{nil}$ **then**
- 11: $\omega_{L'}(\text{succ}_L(v)) \leftarrow \omega_L(v) \oplus \omega_L(\text{succ}_L(v))$
- 12: **end if**
- 13: **end for**
- 14: **for** all $v \notin I$ **do**
- 15: **if** $\text{succ}_L(v) \neq \text{nil}$ **and** $\text{succ}_L(v) \in I$ **then**
- 16: $\text{succ}_{L'}(v) \leftarrow \text{succ}_L(\text{succ}_L(v))$
- 17: **end if**
- 18: **end for**
- 19: Let L' be the list defined by the vertices in $L \setminus I$, pointers $\text{succ}_{L'}(v)$ and weights $\omega_{L'}(v)$.
- 20: Recursively apply procedure FASTLISTRANKING to list L' . Let $\rho_{L'}(v)$ be the rank assigned to every element v in $L \setminus I$.
- 21: **for** all $v \notin I$ **do**
- 22: $\rho_L(v) \leftarrow \rho_{L'}(v)$
- 23: **if** $\text{succ}_L(v) \neq \text{nil}$ **and** $\text{succ}_L(v) \in I$ **then**
- 24: $\rho_L(\text{succ}_L(v)) \leftarrow \rho_L(v) \oplus \omega_L(\text{succ}_L(v))$
- 25: **end if**
- 26: **end for**


Listranking




- Line 4: $O(\text{sort}(N))$



Listranking




- Line 4: $O(\text{sort}(N))$



Listranking

- Line 5-8: $O(\text{scan}(N))$



L'




I




Listranking

- Line 5-8: $O(\text{scan}(N))$




Listranking


- Line 9-13: $O(\text{sort}(N))$



L'




I




Listranking

- Line 14-18: $O(\text{sort}(N))$



L'




I




- Line 19-20: $O(I(2/3N))$
- Line 21-26: $O(\text{sort}(N))$
 - Sort L' based on their weight ranks
 - Sort I based on the weight ranks of their successors


- Replace $\{v, w\}$ with directed edges (v, w) and (w, v)
- $\forall v \in T$:
 - Let incoming edges be e_1, \dots, e_k and outgoing edges be e'_1, \dots, e'_k where e_i and e'_i have the same endpoints
 - edge e_i is succeeded by edge $e'_{i \bmod k}$



- Replace $\{v, w\}$ with directed edges (v, w) and (w, v)
- $\forall v \in T$:
 - Let incoming edges be e_1, \dots, e_k and outgoing edges be e'_1, \dots, e'_k where e_i and e'_i have the same endpoints
 - edge e_i is succeeded by edge $e'_{i \bmod k}$




- Replace $\{v, w\}$ with directed edges (v, w) and (w, v)
- $\forall v \in T$:
 - Let incoming edges be e_1, \dots, e_k and outgoing edges be e'_1, \dots, e'_k where e_i and e'_i have the same endpoints
 - edge e_i is succeeded by edge $e'_{i \bmod k}$




Euler Tour

- Replace $\{v, w\}$ with directed edges (v, w) and (w, v)
- $\forall v \in T$:
 - Let incoming edges be e_1, \dots, e_k and outgoing edges be e'_1, \dots, e'_k where e_i and e'_i have the same endpoints
 - edge e_i is succeeded by edge $e'_{i \bmod k}$



Euler Tour

- Adjacency list representation
 - Euler tour: $O(\text{scan}(N))$
- Unordered collection of edges
 - Euler tour: $O(\text{sort}(N))$



Rooting a tree

- A tree can be rooted in $O(\text{sort}(N))$ I/Os
 - 1: Compute an Euler tour L of tree T
 - 2: Compute the rank of every edges e in L
 - 3: for every edges $\{u, w\} \in T$ do
 - 4: Store the ranks of edges (v, w) and (w, v) in L with $\{u, w\}$

- Labeling
 - Preorder
 - Postorder
 - Depth
- Procedure LabelTree
 - 1: Compute an Euler tour L of tree T that start at the root of T
 - 2: Assign appropriate weights to the edges in the Euler tour
 - 3: Compute the weighted rank of each edges in L
 - 4: Extract a labeling of the vertices of T from these ranks

- Depth

$$w(e) = \begin{cases} 1 & \text{if } v = p(w) \\ -1 & \text{if } w = p(v) \end{cases}$$

- Preorder

$$w(e) = \begin{cases} 1 & \text{if } v = p(w) \\ 0 & \text{if } w = p(v) \end{cases}$$

Evaluating Directed Acyclic Graphs

- Given a DAG $G = (V, E)$
 - Each vertex is associated with $w(v)$ (known) and $(\rho(v))$ (to be computed)
 - $\rho(v)$ depends on the in-neighbors u_1, \dots, u_k of v
 - Listranking is a special case
 - Two assumptions to get efficient solution
 - Vertices are given in a topological sort, otherwise $\Omega(|V|)$ I/Os are needed to topologically sort vertices
 - If the in-degree is unbounded, computation of $\rho(v)$ from its in-neighbors u_1, \dots, u_k can be done in $O(\text{sort}(k))$ I/Os
- * Since Listranking is so restricted without two above assumptions we get efficient solution

Time-Forward Processing

- Procedure TimeForwardProcessing

1: $Q \leftarrow \emptyset$

2: For every vertex $v \in G$ in topologically sorted order do

3: Let u_1, \dots, u_k be in-neighbors of v

4: Retrieve $\rho(u_1), \dots, \rho(u_k)$ from Q using k DeleteMin operations

4: Compute $\rho(v)$ from $w(v)$ and $\rho(u_1), \dots, \rho(u_k)$

5: Let w_1, \dots, w_ℓ be out-neighbors of v

6: Insert ℓ copies of $\rho(v)$ into priority queue Q . Give the i -th copy priority w_i

- A DAG G can be evaluated in $O(\text{sort}(E))$ I/Os if vertices are given a topologically sorted order

Maximal Independent Set

- Procedure `MaximalIndependentSet`

- 1: $I \leftarrow \emptyset$

- 2: Direct the edge of G from vertices with lower numbers to vertices with higher numbers

- 3: Sort the vertices of G by their numbers and the edges by the number of their sources

- 4: for every vertices $v \in G$ in sorted order

- 4: if no in-neighbor of v is in I then

- 5: add v to I

- Line 4-8 can be simulated using `Time-Forward Processing`

- A maximal independent set of a undirected graph G can be computed in $O(\text{sort}(|V| + |E|))$

Maximal Independent Set

- Any maximal independent set of a list L has size at least $N/3$, since every vertex has at most two neighbors
- A maximal independent set of a list L can be computed in $O(\text{sort}(N))$

- Parallel Random Access Machine (PRAM)
 - N processors
 - Shared Memory
- Read/write conflicts
 - Exclusive Read Exclusive Write (EREW)
 - Concurrent Read Exclusive Write (CREW)
 - Exclusive Read Concurrent Write (ERCW)
 - Concurrent Read Concurrent Write (CRCW)



- Assumptions

- N processors and N space
- EREW strategy
- In a single step, each PRAM processor reads $O(1)$ operands from memory, performs some computation, and then writes $O(1)$ results to memory.

- Simulation

- Sort a copy of the contents of the PRAM memory based on the indices of the processors for which they will be operands in this step.
- Scan this copy and perform the computation for each processor being simulated, and write the results to the disk as we do so
- Sort the results of the computation based on the memory addresses to which the PRAM processors would store them and then scan the list and a reserved copy of memory to merge the stored values back into the memory.

PRAM Simulation



- If a PRAM algorithm using $O(N)$ space and processors runs in T steps, the algorithm can be simulated using $O(T \cdot \text{sort}(N))$ I/Os
- If every $O(1)$ steps, space and the number of processors decrease by a constant factor of N , the algorithm can be simulated in $O(\text{sort}(N))$ I/Os.

Summary: Algorithms for trees

- **Listranking** can be performed in $O(\text{sort}(N))$ I/Os
- The following algorithms can be done on trees using Listranking
 - Making rooted
 - Preorder ranking
 - Postorder ranking
 - Computing depth
- **Techniques**
 - Time-forward processing
 - PRAM simulation

- **I/O efficient graph algorithms**

Lecture notes by Norbert Zeh.

- Section 1-4