


Massive Data Algorithmics


Lecture 3: External Search Trees

Binary search tree

- Standard method for search among N elements
- We assume elements in leaves
- Search traces at least one root-leaf path
- If nodes stored arbitrarily on disk
 - \Rightarrow Search in $O(\log_2 N)$ I/Os
 - \Rightarrow Range-search in $O(\log_2 N + T)$ I/Os




BFS Blocking



- Block height: $O(\log_2 N) / O(\log_2 B) = O(\log_B N)$
- Output elements blocked \Rightarrow Range-search in $O(\log_B N + T/B)$ I/Os
- Optimal:** $O(N/B)$ space and $O(\log_B N + T/B)$ query

Updating


- Maintaining BFS blocking during updates?
 - Balance normally maintained in search trees using rotations




- Seems very difficult to maintain BFS blocking during rotation
 - Also need to make sure output (leaves) is blocked!

B-trees

- BFS-blocking naturally corresponds to tree with fan-out $\theta(B)$




- B-trees balanced by allowing node degree to vary
 - Re-balancing performed by splitting and merging nodes



(a, b) -Trees


- T is an (a, b) -tree ($a \geq 2$ and $b \geq 2a - 1$)
 - All leaves on the same level and contain between a and b elements
 - Except for the root, all nodes have degree between a and b
 - Root has degree between 2 and b



- (a, b) -tree uses linear space and has height $O(\log_a N)$
- Choosing $a, b = \Theta(B)$, each node/leaf stored in one disk block
- $O(N/B)$ space and $O(\log_B N + T/B)$ query


(a, b) -Trees Insert

- Search and insert element in leaf v
- DO v has $b+1$ elements/children
 - make nodes v' and v'' with $\lfloor (b+1)/2 \rfloor$ and $\lceil (b+1)/2 \rceil$ elements
 - insert element (ref) in $\text{parent}(v)$
(make new root if necessary)
- $v = \text{parent}(v)$





Insert touch $O(\log_a N)$ nodes


Example: (2,4)-Tree Insert



Example: (2,4)-Tree Insert









(a, b) -Trees Deletion

- Search and delete element from leaf v
- DO v has $a - 1$ elements/children
 - Fuse v with sibling v' :
 - move children of v' to v
 - delete element (ref) from parent(v)
(delete root if necessary)
 - If v has $> b$ (and $\leq a + b - 1 < 2b$) children
split v
- $v = \text{parent}(v)$




Delete touch $O(\log_a N)$ nodes


Example: (2,4)-Tree Delete




Example: (2,4)-Tree Delete




Example: (2,4)-Tree Delete



Example: (2,4)-Tree Delete




Example: (2,4)-Tree Delete



(a, b) -Trees Properties

- If $b = 2a - 1$ every update can cause many re-balancing operations



- If $b \geq 2a$ update only cause $O(1)$ re-balancing operations amortized
- If $b > 2a$ only $O(1/(b/2 - a)) = O(1/a)$ re-balancing operations amortized
 - *Both somewhat hard to show
- If $b=4a$ easy to show that update causes $O(1/a \log_a N)$ re-balance operations amortized
 - * After split during insert a leaf contains $\cong 4a/2 = 2a$ elements
 - * After fuse during delete a leaf contains between $\cong 2a$ and $\cong 5a$ elements (split if more than $3a \Rightarrow$ between $3/2a$ and $5/2a$)

Summary and Conclusion: B-trees

- B-trees: (a, b) -trees with $a, b = \Theta(B)$
 - $O(N/B)$ space
 - $O(\log_B N + T/B)$ query
 - $O(\log_B N)$ update
- B-trees with elements in the leaves sometimes called **B⁺-trees**
- Construction in $O(N/B \log_{M/B} N/B)$ I/Os
 - Sort elements and construct leaves
 - Build tree level-by-level bottom-up

Summary and Conclusion: B-trees

- B-tree with branching parameter b and leaf parameter k ($b, k \geq 8$)
 - All leaves on same level and contain between $1/4k$ and k elements
 - Except for the root, all nodes have degree between $1/4b$ and b
 - Root has degree between 2 and b
- B-tree with leaf parameter $k = \Omega(B)$
 - $O(N/B)$ space
 - Height $O(\log_b N/B)$
 - $O(1/k)$ amortized leaf rebalance operations
 - $O(1/(bk) \log_b N/B)$ amortized internal node rebalance operations
- B-tree with branching parameter B^c , $0 < c \leq 1$, and leaf parameter B
 - Space $O(N/B)$, updates $O(\log_B N)$, queries $O(\log_B N + T/B)$

References

- **External Memory Geometric Data Structures**

Lecture notes by Lars Arge.

- Section 1-3