


Massive Data Algorithmics

Lecture 7: Range Searching


Three-Sided Range Queries

- Interval management: 1.5 dimensional search



- More general $2d$ problem: Dynamic 3-sided range searching


- Maintain set of points in plane such that given query (q_1, q_2, q_3) , all points (x, y) with $q_1 \leq x \leq q_2$ and $y \geq q_3$ can be found efficiently



Three-Sided Range Queries

- Static solution:

- Sweep top-down inserting x in persistent B-tree at (x, y)
- Answer query by performing range query with $[q_1, q_2]$ in B-tree at q_3




- Optimal:

- $O(N/B)$ space
- $O(\log_B N + T/B)$ query
- $O(N/B \log_{M/B} N/B)$ construction


- **Dynamic?** in internal memory: priority search tree

Internal Priority Search Tree




- Base tree on x-coordinates with nodes augmented with points
- Heap on y-coordinates:
 - Decreasing y values on root-leaf path
 - (x, y) on path from root to leaf holding x
 - If v holds point then $\text{parent}(v)$ holds point

Internal Priority Search Tree




- Linear space
 - Insert of (x, y) (assuming fixed x -coordinate set):
 - Compare y with y -coordinate in root
 - Smaller: Recursively insert (x, y) in subtree on path to x
 - Bigger: Insert in root and recursively insert old point in subtree
- $\Rightarrow O(\log N)$ update

Internal Priority Search Tree



- Linear space
 - Insert of (x, y) (assuming fixed x -coordinate set):
 - Compare y with y -coordinate in root
 - Smaller: Recursively insert (x, y) in subtree on path to x
 - Bigger: Insert in root and recursively insert old point in subtree
- $\Rightarrow O(\log N)$ update


Internal Priority Search Tree



- Linear space
- Insert of (x,y) (assuming fixed x -coordinate set):
 - Compare y with y -coordinate in root
 - Smaller: Recursively insert (x,y) in subtree on path to x
 - Bigger: Insert in root and recursively insert old point in subtree

$\Rightarrow O(\log N)$ update


Internal Priority Search Tree



- Linear space
- Insert of (x,y) (assuming fixed x -coordinate set):
 - Compare y with y -coordinate in root
 - Smaller: Recursively insert (x,y) in subtree on path to x
 - Bigger: Insert in root and recursively insert old point in subtree

$\Rightarrow O(\log N)$ update


Internal Priority Search Tree



- Linear space
- Insert of (x,y) (assuming fixed x -coordinate set):
 - Compare y with y -coordinate in root
 - Smaller: Recursively insert (x,y) in subtree on path to x
 - Bigger: Insert in root and recursively insert old point in subtree


$\Rightarrow O(\log N)$ update

Internal Priority Search Tree




- Query with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2
- $\Rightarrow O(\log N + T)$ query

Internal Priority Search Tree




- Query with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2
- $\Rightarrow O(\log N + T)$ query

Internal Priority Search Tree




- Query with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2
- $\Rightarrow O(\log N + T)$ query

Internal Priority Search Tree




- Query with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2
- $\Rightarrow O(\log N + T)$ query

Internal Priority Search Tree




- Query with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2
- $\Rightarrow O(\log N + T)$ query

Internal Priority Search Tree



- Query with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2
- $\Rightarrow O(\log N + T)$ query

Externalizing Priority Search Tree




• Natural idea: Block tree

• Problem:

- $O(\log_B N)$ I/Os to follow paths to q_1 and q_2
- But $O(T)$ I/Os may be used to visit other nodes ("overshooting")

$\Rightarrow O(\log_B N + T)$ query

Externalizing Priority Search Tree




● Solution idea:

- Store B points in each node:
 - * $O(B^2)$ points stored in each supernode
 - * B output points can pay for **overshooting**
- Bootstrapping:
 - * Store $O(B^2)$ points in each supernode in static structure


Externalizing Priority Search Tree

- **Base tree:** Weight-balanced B-tree with branching parameter $B/4$ and leaf parameter B on x -coordinates
- Points in **heap order**:
 - Root stores B top points for each of the $\Theta(B)$ child slabs
 - Remaining points stored recursively
- Points in each node stored in B^2 -structure
 - Persistent B-tree structure for static problem
- \Rightarrow Linear space




Externalizing Priority Search Tree

- Query with (q_1, q_2, q_3) starting at root v :
 - Query B^2 -structure and report points satisfying query
 - Visit child v if
 - * v on path to q_1 or q_2
 - * All points corresponding to v satisfy query




Externalizing Priority Search Tree

- Query with (q_1, q_2, q_3) starting at root v :
 - Query B^2 -structure and report points satisfying query
 - Visit child v if
 - * v on path to q_1 or q_2
 - * All points corresponding to v satisfy query



Externalizing Priority Search Tree

- Query with (q_1, q_2, q_3) starting at root v :
 - Query B^2 -structure and report points satisfying query
 - Visit child v if
 - * v on path to q_1 or q_2
 - * All points corresponding to v satisfy query




Externalizing Priority Search Tree

Analysis:


- $O(\log_B B^2 + T_v/B) = O(1 + T_v/B)$ I/Os used to visit node v
- $O(\log_B N)$ nodes on path to q_1 or q_2
- For each node v not on path to q_1 or q_2 visited, B points reported in $\text{parent}(v)$

$$\Rightarrow O(\log_B N + T/B)$$




Externalizing Priority Search Tree

- **Insert** (x, y) (ignoring insert in base tree - rebalancing):
 - Find relevant node u :
 - * Query B^2 -structure to find B points in root corresponding to node u on path to x
 - * If y smaller than y -coordinates of all B points then recursively search in u
 - Insert (x, y) in B^2 -structure of v
 - If B^2 -structure contains $> B$ points for child u , remove lowest point and insert recursively in u
- **Delete**: Similarly




Externalizing Priority Search Tree

- **Insert** (x,y) (ignoring insert in base tree - rebalancing):
 - Find relevant node u :
 - * Query B^2 -structure to find B points in root corresponding to node u on path to x
 - * If y smaller than y -coordinates of all B points then recursively search in u
 - Insert (x,y) in B^2 -structure of v
 - If B^2 -structure contains $> B$ points for child u , remove lowest point and insert recursively in u
- **Delete**: Similarly




Externalizing Priority Search Tree

- **Insert** (x,y) (ignoring insert in base tree - rebalancing):
 - Find relevant node u :
 - * Query B^2 -structure to find B points in root corresponding to node u on path to x
 - * If y smaller than y -coordinates of all B points then recursively search in u
 - Insert (x,y) in B^2 -structure of v
 - If B^2 -structure contains $> B$ points for child u , remove lowest point and insert recursively in u
- **Delete**: Similarly




Externalizing Priority Search Tree

- **Insert** (x,y) (ignoring insert in base tree - rebalancing):
 - Find relevant node u :
 - * Query B^2 -structure to find B points in root corresponding to node u on path to x
 - * If y smaller than y -coordinates of all B points then recursively search in u
 - Insert (x,y) in B^2 -structure of v
 - If B^2 -structure contains $> B$ points for child u , remove lowest point and insert recursively in u
- **Delete**: Similarly



Externalizing Priority Search Tree

- **Insert** (x,y) (ignoring insert in base tree - rebalancing):
 - Find relevant node u :
 - * Query B^2 -structure to find B points in root corresponding to node u on path to x
 - * If y smaller than y -coordinates of all B points then recursively search in u
 - Insert (x,y) in B^2 -structure of v
 - If B^2 -structure contains $> B$ points for child u , remove lowest point and insert recursively in u
- **Delete**: Similarly



Externalizing Priority Search Tree

● Analysis:

- Update visits $O(\log_B N)$ nodes
- B^2 -structure queried/updated in each node
 - * One query
 - * One insert and one delete

● B^2 -structure analysis:

- Query: $O(\log_B B^2 + B/B) = O(1)$ I/Os
- Update: $O(1)$ using global rebuilding
 - * Store updates in update block
 - * Rebuild after B updates using $O(B^2/B \log_{M/B} BB^2/B) = O(B)$ I/Os

$\Rightarrow O(\log_B N)$ update

Dynamic Base Tree

Deletion:


- Delete point as previously
- Delete x -coordinate from base tree using global rebuilding

$\Rightarrow O(\log_B N)$ I/Os

Insertion:

- Insert x -coordinate in base tree and rebalance (using splits)
- Insert point as previously

- **Split:** Boundary in v becomes boundary in $\text{parent}(v)$



Dynamic Base Tree

- **Split:** When v splits B new points needed in $\text{parent}(v)$
- One point obtained from v' (v'') using bubble-up operation:


- Find top point p in v'
- Insert p in B^2 -structure
- Remove p from B^2 -structure of v'
- Recursively bubble-up point to v'

$\Rightarrow O(\log_B N)$ I/Os

- **Bubble-up** in $O(\log_B w(v))$ I/Os

- Follow one path from v to leaf
- Uses $O(1)$ I/O in each node

\Rightarrow Split in $O(B \log_B w(v)) = O(w(v))$ I/Os



Dynamic Base Tree

- $O(1)$ amortized split cost:
 - Cost: $O(w(v))$
 - Weight balanced base tree: $\Omega(w(v))$ inserts below v between splits





- External Priority Search Tree
 - Space: $O(N/B)$
 - Query: $O(\log_B N + T/B)$
 - Update: $O(\log_B N)$ I/Os amortized

Summary/Conclusion: Range Search

- We have now discussed structures for special cases of two-dimensional range searching

- Space: $O(N/B)$
- Query: $O(\log_B N + T/B)$
- Updates: $O(\log_B N)$



- Cannot be obtained for general (4-sided) 2d range searching:

- $O(\log_B^c N)$ query requires $\Omega(\frac{N}{B} \frac{\log_B N}{\log_B \log_B N})$ space
- $O(N/B)$ space requires $\Omega(\sqrt{N/B})$ query

